

1.	Start MongoDB
1.1	<p>start up the vagrant box</p> <ul style="list-style-type: none"> - navigate to the redis/vagrant folder: <code>cd mongodb/vagrant</code> - enter the command: <code>vagrant up</code>
This will now download and install the vagrant box, this might take a while	
1.2	<p>start up the mongodb shell.</p> <ul style="list-style-type: none"> - enter the command: <code>vagrant ssh</code> - you will now go the the vagrant box - correct a local setting via: <code>export LC_ALL=C</code> - start the shell via entering: <code>mongo</code> <p>You should see <code>></code></p> <p>If you see an error: Failed global initialization: BadValue Invalid or no user locale set. Please ensure LANG and/or LC_* environment variables are set correctly</p> <p>Then first run: <code>export LC_ALL=C</code></p> <p>Followed by <code>mongo</code></p>
1.3	<p>let us first switch to the test database via the command: <code>use test</code></p> <p>this will represent the message: <code>switched to db test</code></p> <p>If a database does not exist - mongodb will automatically create it - it is very flexible.</p>
2.	Insert a document
2.1	<p>Insert a document into a collection restaurants</p> <pre>db.restaurants.insert({ "address" : { "street" : "2 Avenue", "zipcode" : "10075",</pre>

```

        "building" : "1480",
        "coord" : [ -73.9557413, 40.7720266 ],
    },
    "borough" : "Manhattan",
    "cuisine" : "Italian",
    "grades" : [
        {
            "date" : ISODate("2014-10-01T00:00:00Z"),
            "grade" : "A",
            "score" : 11
        },
        {
            "date" : ISODate("2014-01-16T00:00:00Z"),
            "grade" : "B",
            "score" : 17
        }
    ],
    "name" : "Vella",
    "restaurant_id" : "41704620"
}
)

```

The method returns a WriteResult object with the status of the operation:
WriteResult({ "nInserted" : 1 })

If the _id field is not passed in as an argument, mongodb will automatically add it.

2.2

To find documents, just use the find command:

```

db.restaurants.find()
db.restaurants.find( { "borough": "Manhattan" } )
db.restaurants.find( { "address.zipcode": "10075" } )

```

Via the dot annotation you can query on properties of properties

```

db.restaurants.find( { "grades.grade": "B" } )

```

	<p>At the .pretty() function to make the output more readable</p> <pre>db.restaurants.find({ "grades.grade": "B" }).pretty()</pre>
2.3	<p>The greater then operator \$gt</p> <pre>db.restaurants.find({ "grades.score": { \$gt: 3 } })</pre> <p>You also have a less then \$lt operator</p>
2.4	<p>AND clause</p> <p>Just separate the properties by a , in the query:</p> <pre>db.restaurants.find({ "cuisine": "Italian", "address.zipcode": "10075" })</pre>
2.5	<p>Logical OR</p> <p>Via the \$or operator you can specify an OR clause</p> <pre>db.restaurants.find({ \$or: [{ "cuisine": "Italian" }, { "address.zipcode": "10075" }] })</pre>
2.6	<p>Sort query results</p> <p>To specify an order for the result set append the sort() method to the query. Pass the fields to sort by and the sort type (1 for ascending and -1 for descending)</p> <pre>db.restaurants.find().sort({ "borough": 1, "address.zipcode": 1 })</pre>
2.7	<p>Update fields:</p> <p>The following operation updates the first document with name equal to "Juni", using the \$set operator to update the cuisine field and the \$currentDate operator to update the lastModified field with the current date.</p> <pre>db.restaurants.update({ "name": "Juni" }, { \$set: { "cuisine": "American (New)" }, \$currentDate: { "lastModified": true } })</pre>

	<pre>)</pre> <p>Embedded fields can be updated via the “dot” operator:</p> <pre>db.restaurants.update({ "restaurant_id" : "41156888" }, { \$set: { "address.street": "East 31st Street" } })</pre>
2.8	<p>By default update only updates a single document.</p> <p>In order to update more than 1 document the “multi” option must be set to true</p> <pre>db.restaurants.update({ "address.zipcode": "10016", cuisine: "Other" }, { \$set: { cuisine: "Category To Be Determined" }, \$currentDate: { "lastModified": true } }, { multi: true })</pre>
2.9	<p>To replace an entire document by another => pass an entirely new document as 2nd parameter in the update method.</p> <p>Important: after the update the document will only contain the fields of the provided document and the id.</p> <pre>db.restaurants.update({ "restaurant_id" : "41704620" }, { "name" : "Vella 2", "address" : { "coord" : [-73.9557413, 40.7720266], "building" : "1480", "street" : "2 Avenue", "zipcode" : "10075" } })</pre> <p>You can check this via</p> <pre>db.restaurants.find({ "name" : "Vella 2" })</pre>

2.10	<p>Remove documents via remove.</p> <pre>db.restaurants.remove({ "borough": "Manhattan" })</pre> <p>By default the remove function removes all documents matching to the condition.</p> <p>You can pass the justOne option to make sure that only 1 documents gets deleted:</p> <pre>db.restaurants.remove({ "borough": "Queens" }, { justOne: true })</pre>
2.11	<p>With an empty remove you can remove all documents of a collection:</p> <pre>db.restaurants.remove()</pre>
2.12	<p>Via drop you can drop a collection:</p> <pre>db.restaurants.drop()</pre> <p>If the collection to drop does not exist, the operation will return false.</p>
3.	<h2>Aggregate queries</h2>
3.1	<p>Load in the complete test dataset:</p> <p>Open a new vagrant ssh shell:</p> <ul style="list-style-type: none"> - Open a command prompt - Navigate to the vagrant folder - Enter "vagrant ssh" - Go to the vagrant folder via cd vagrant - execute the following command: <pre>mongoimport --db test --collection restaurants --drop --file primer-dataset.json</pre>
3.2	<p>Group by a field and calculate a count</p> <p>Via \$group you can group by a specified key.</p> <p>In the \$group stage, specify the group by key in the _id field.</p> <p>\$group accesses fields by the field path, which is the field name prefixed by \$</p> <p>The \$group stage can use accumulators to perform calculations on each group</p> <p>Over here we will use the \$sum accumulator</p>

	<pre>db.restaurants.aggregate([{ \$group: { "_id": "\$borough", "count": { \$sum: 1 } } }]);</pre> <p>The results consists of the following documents:</p> <pre>{ "_id" : "Staten Island", "count" : 969 } { "_id" : "Brooklyn", "count" : 6086 } { "_id" : "Manhattan", "count" : 10259 } { "_id" : "Queens", "count" : 5656 } { "_id" : "Bronx", "count" : 2338 } { "_id" : "Missing", "count" : 51 }</pre>
3.3	<p>Use the \$match to filter documents. \$match uses the query syntax.</p> <pre>db.restaurants.aggregate([{ \$match: { "borough": "Queens", "cuisine": "Brazilian" } }, { \$group: { "_id": "\$address.zipcode", "count": { \$sum: 1 } } }]);</pre> <p>The result will be</p> <pre>{ "_id" : "11368", "count" : 1 } { "_id" : "11106", "count" : 3 } { "_id" : "11377", "count" : 1 } { "_id" : "11103", "count" : 1 } { "_id" : "11101", "count" : 2 }</pre> <p>Where _id contains the zipcode value - which was used to group by.</p>
3.4	<p>There exist some common aggregation methods.</p> <p>count:</p> <p>db.restaurants.count() will return the total number of restaurant documents in the collection</p> <p>count also works after a find():</p> <pre>db.restaurants.find({ "cuisine": "Italian", "address.zipcode": "10075" }).count() => 15</pre>

	<p>distinct:</p> <p>The distinct() function allows you to retrieve the distinct values of a single field.</p> <pre>db.restaurants.distinct("cuisine")</pre>
Indexes	
4.1	<p>Create a single field index</p> <p>Create an ascending index on the cuisine field of the restaurants collection</p> <pre>db.restaurants.createIndex({ "cuisine": 1 })</pre> <p>This returns:</p> <pre>{ "createdCollectionAutomatically" : false, "numIndexesBefore" : 1, "numIndexesAfter" : 2, "ok" : 1 }</pre>
4.2	<p>Compound indexes are also possible</p> <pre>db.restaurants.createIndex({ "cuisine": 1, "address.zipcode": -1 })</pre>
5. Map Reduce	
	<p>MongoDB can use the map - reduce type of working to handle more complex tasks</p>
5.1	<p>Data setup:</p> <pre>db.sessions.save({ userid: "a", ts: ISODate('2011-11-03 14:17:00'), length: 95 }); db.sessions.save({ userid: "b", ts: ISODate('2011-11-03 14:23:00'), length: 110 }); db.sessions.save({ userid: "c", ts: ISODate('2011-11-03 15:02:00'), length: 120 }); db.sessions.save({ userid: "d", ts: ISODate('2011-11-03 16:45:00'), length: 45 }); db.sessions.save({ userid: "a", ts: ISODate('2011-11-04 11:05:00'), length: 105 }); db.sessions.save({ userid: "b", ts: ISODate('2011-11-04 13:14:00'), length: 120 });</pre>

	<pre>db.sessions.save({ userid: "c", ts: ISODate('2011-11-04 17:00:00'), length: 130 }); db.sessions.save({ userid: "d", ts: ISODate('2011-11-04 15:37:00'), length: 65 });</pre>
5.2	<p>Define a map function to map the userid to an object that contains the fields userid, total_time, count and avg_time</p> <pre>var mapFunction = function() { var key = this.userid; var value = { userid: this.userid, total_time: this.length, count: 1, avg_time: 0 }; emit(key, value); };</pre>
	<p>Define the corresponding reduce function with 2 arguments, keys and values to calculate the total time and the count. Key = the userid while values is an array with properties.</p> <pre>var reduceFunction = function(key, values) { var reducedObject = { userid: key, total_time: 0, count: 0, avg_time: 0 }; values.forEach(function(value) { reducedObject.total_time += value.total_time; reducedObject.count += value.count; }); return reducedObject; };</pre>
	<p>Define the finalize function with two arguments key and reducedValue. The function modifies the reducedValue document to add another field average and returns the modified document.</p> <pre>var finalizeFunction = function (key, reducedValue) {</pre>

	<pre> if (reducedValue.count > 0) reducedValue.avg_time = reducedValue.total_time / reducedValue.count; return reducedValue; };</pre>
	<p>Now perform the map reduce operation.</p> <p>The results will be written out to the collection "session_stat" If that collection already exists, it will be emptied and overwritten</p> <pre>db.sessions.mapReduce(mapFunction, reduceFunction, { out: "session_stat", finalize: finalizeFunction })</pre> <p>Via <code>db.session_stat.find().pretty()</code> you can see its contents</p>